

HomeWork 7

This homework need to use PEFT(Parameter Efficient Fine-Tuning) to fine-tune a LLM(Large Language Model) and try to use different LoRA configs finally compare the result.

Setup

```
!pip install transformers
!pip install datasets
!pip install accelerate
!pip install peft
```

Result

Without LoRA

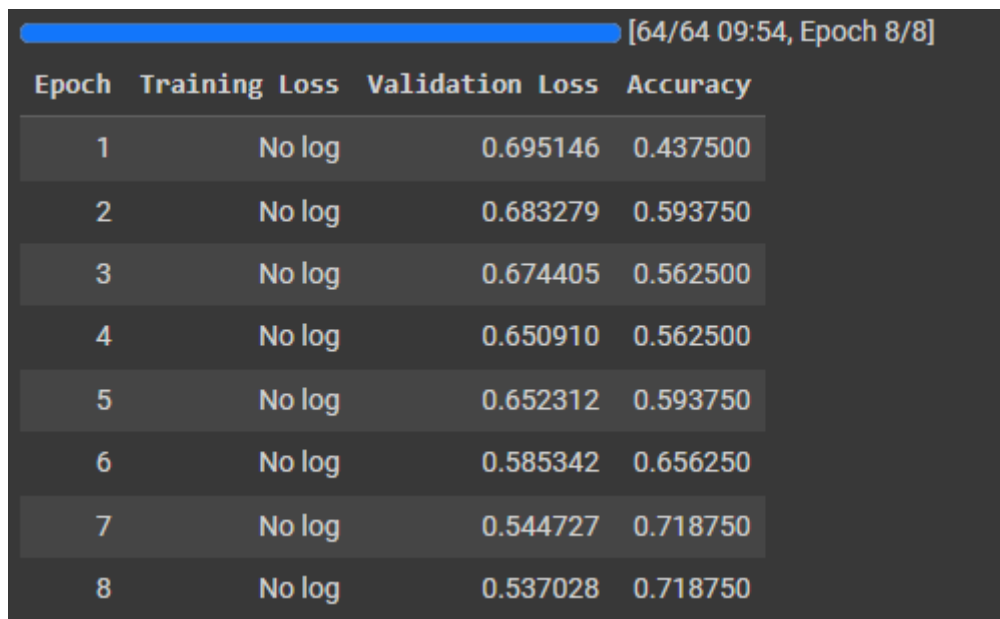
trainable model parameters: 65783042 all model parameters: 65783042 percentage of trainable model parameters: 100.00%

Using LoRA

trainable model parameters: 813314 all model parameters: 66596356 percentage of trainable model parameters: 1.22%

Config 1

```
r=8,
lora_alpha=32,
lora_dropout=0.05,
target_modules=["q_lin", "k_lin", "v_lin"],
bias='none',
task_type=TaskType.SEQ_CLS
```




[64/64 09:54, Epoch 8/8]

Epoch	Training Loss	Validation Loss	Accuracy
1	No log	0.695146	0.437500
2	No log	0.683279	0.593750
3	No log	0.674405	0.562500
4	No log	0.650910	0.562500
5	No log	0.652312	0.593750
6	No log	0.585342	0.656250
7	No log	0.544727	0.718750
8	No log	0.537028	0.718750

```
PredictionOutput(predictions=array([[0.20563483, 0.21584545],
[0.22710314, 0.19645725],
[0.19117099, 0.14843869],
[0.21138549, 0.16851144],
[0.15074639, 0.20138836],
[0.21285076, 0.16002882],
[0.20809 , 0.11896151],
[0.21061125, 0.19299033],
[0.17564015, 0.23414241],
[0.17190261, 0.19909424],
[0.19763005, 0.2163772 ],
[0.2524781 , 0.10385972],
[0.18600446, 0.16466182],
[0.19018903, 0.16652927],
[0.19784276, 0.16024332],
[0.17829013, 0.10582904],
[0.22059232, 0.12387354],
[0.18845229, 0.17890874],
[0.21097481, 0.14440343],
[0.14287387, 0.20949566],
[0.20877954, 0.14217615],
[0.22035219, 0.18926692],
[0.17945167, 0.18837884],
[0.21406782, 0.18342996],
[0.22085267, 0.178282 ],
[0.19710688, 0.18896426],
[0.1706594 , 0.21593145],
[0.23764741, 0.14943534],
[0.17323879, 0.17553255],
[0.17276055, 0.12632418],
[0.19190821, 0.19915201],
[0.17582431, 0.2054169 ]], dtype=float32), label_ids=array([[1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0,
0, 1, 1, 0, 1, 1, 1, 0, 0, 1]), metrics={'test_loss': 0.6861392855644226, 'test_accuracy': 0.65025, 'test_runtime': 4.6837, 'test_samples_per_second': 6.832, 'test_steps_per_second': 0.427})
```

Config 2

r=8,
 lora_alpha=16,
 lora_dropout=0.01,
 target_modules=["q_lin", "k_lin", "v_lin"],
 bias='none',
 task_type=TaskType.SEQ_CLS



[64/64 11:37, Epoch 8/8]

Epoch	Training Loss	Validation Loss	Accuracy
1	No log	0.688975	0.625000
2	No log	0.667406	0.687500
3	No log	0.631328	0.750000
4	No log	0.551050	0.843750
5	No log	0.477940	0.781250
6	No log	0.432638	0.843750
7	No log	0.400708	0.843750
8	No log	0.393947	0.843750

```
PredictionOutput(predictions=array([[ -0.42432398, 0.81219757],
[ 0.7792884 , -0.2631061 ],
[ -0.29699902, 0.3365784 ],
[ -0.11130465, 0.32808650],
[ -0.95701617, 0.99197636],
[ 0.60232978, -0.33735508],
[ 0.7076106 , -0.8032887 ],
[ 0.81153014, -0.7631841 ],
[ -0.94923956, 0.98996325],
[ -1.1813029 , 0.92249143],
[ -0.69220376, 0.910442 ],
[ 1.1302046 , -1.1446705 ],
[ 0.7094246 , -0.44799092],
[ 0.09983131, 0.03113251],
[ 0.8099524 , -0.6487731 ],
[ 0.9314997 , -1.0685852 ],
[ -0.13552408, 0.06428724],
[ -1.0934927 , 0.90159667],
[ -0.3074137 , 0.29185358],
[ -1.1500902 , 0.9405169 ],
[ 0.47950438, -0.3754086 ],
[ 0.22833812, -0.05777533],
[ -0.0194078 , 0.07729631],
[ -0.72742295, 0.8095778 ],
[ 0.54920965, -0.3115983 ],
[ -0.62890846, 0.5672336 ],
[ -0.7918495 , 1.0297548 ],
[ -0.42977846, 0.34765667],
[ 0.3958242 , -0.3400324 ],
[ 0.2913689 , -0.2709194 ],
[ 0.57326937, -0.16788958],
[ -1.1254654 , 1.0232337 ]], dtype=float32), label_ids=array([[1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 0,
0, 1, 1, 0, 1, 1, 1, 0, 0, 1]), metrics={'test_loss': 0.3939467668533325, 'test_accuracy': 0.84375, 'test_runtime': 12.2686, 'test_samples_per_second': 2.604, 'test_steps_per_second': 0.163})
```

Config 3

```
r=16,  
lora_alpha=8,  
lora_dropout=0.002,  
target_modules=["q_lin", "k_lin","v_lin"],  
bias='none',  
task_type=TaskType.SEQ_CLS
```

[64/64 10:43, Epoch 8/8]

Epoch	Training Loss	Validation Loss	Accuracy
1	No log	0.688975	0.625000
2	No log	0.667406	0.687500
3	No log	0.631328	0.750000
4	No log	0.551050	0.843750
5	No log	0.477940	0.781250
6	No log	0.432638	0.843750
7	No log	0.400708	0.843750
8	No log	0.393947	0.843750

```
PredictionOutput(predictions=array([[ -0.42432398,  0.81219757],  
[ 0.7792884 , -0.2631061 ],  
[ -0.25695962,  0.3365784 ],  
[ -0.11180465,  0.32808656],  
[ -0.95701617,  0.99187636],  
[ 0.60232973, -0.33735508],  
[ 0.7076106 , -0.8032887 ],  
[ 0.91153014, -0.7831641 ],  
[ -0.94923955,  0.98396325],  
[ -1.1812029 ,  0.92249143],  
[ -0.69220376,  0.910442 ],  
[ 1.1302046 , -1.1446705 ],  
[ 0.7094246 , -0.44799092],  
[ 0.09983131,  0.03113251],  
[ 0.8099524 , -0.6487731 ],  
[ 0.9314997 , -1.0685852 ],  
[ -0.13552408,  0.06428724],  
[ -1.0934927 ,  0.90159667],  
[ -0.3074137 ,  0.29185358],  
[ -1.1500902 ,  0.9405169 ],  
[ 0.47950438, -0.3754096 ],  
[ 0.22833812, -0.05777533],  
[ -0.0194076 ,  0.07729631],  
[ -0.72742295,  0.8095778 ],  
[ 0.54920965, -0.3115983 ],  
[ -0.62890846,  0.5672336 ],  
[ -0.7918495 ,  1.0297548 ],  
[ -0.42977846,  0.34765667],  
[ 0.3958242 , -0.3400324 ],  
[ 0.2913689 , -0.2709194 ],  
[ 0.57326937, -0.16768958],  
[ -1.1254654 ,  1.0232337 ]], dtype=float32), label_ids=array([1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1, 1, 1, 0, 0,  
0, 1, 1, 0, 1, 1, 1, 0, 0, 1]), metrics={'test_loss': 0.3939467668533325, 'test_accuracy': 0.84375, 'test_runtime': 4.7597, 'test_samples_per_second': 6.723, 'test_steps_per_second': 0.42})
```